

```
In [23]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import zscore
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder
from sklearn.decomposition import PCA
import numpy as np
from sklearn.ensemble import RandomForestClassifier
from sklearn.feature_selection import SelectFromModel
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.feature_selection import SelectFromModel
from scipy import stats
from sklearn.decomposition import PCA
from scipy.stats import chi2_contingency
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import classification_report, accuracy_score
from sklearn.preprocessing import PolynomialFeatures
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.model_selection import cross_val_score
from xgboost import XGBClassifier
from sklearn.model_selection import KFold, cross_val_score
from sklearn.metrics import classification_report, roc_auc_score
import warnings
import joblib
```

```
In [24]: # Suppress warnings
warnings.filterwarnings("ignore")
```

```
In [25]: # Load the cleaned dataset
df = pd.read_csv("Network_anomaly_data.csv")

# Check the first few rows of the data
print(df.head())

# Get general information about the dataset
print(df.info())
```

	duration	protocol	type	service	flag	srcbytes	dstbytes	land
0	0	tcp	ftp_data	SF	491	0	0	
1	0	udp	other	SF	146	0	0	
2	0	tcp	private	S0	0	0	0	
3	0	tcp	http	SF	232	8153	0	
4	0	tcp	http	SF	199	420	0	

	wrongfragment	urgent	hot	...	dsthostsamesrvrate	dsthostdif
fsrvrate \						
0	0	0	0	...	0.17	
0.03						
1	0	0	0	...	0.00	
0.60						
2	0	0	0	...	0.10	
0.05						
3	0	0	0	...	1.00	
0.00						
4	0	0	0	...	1.00	
0.00						

	dsthostsamesrcportrate	dsthostsrvidffhostrate	dsthostserrorrate
te \			
0	0.17	0.00	0.
00			
1	0.88	0.00	0.
00			
2	0.00	0.00	1.
00			
3	0.03	0.04	0.
03			
4	0.00	0.00	0.
00			

	dsthostsrvserrorrate	dsthostrerrorrate	dsthostsrvrerrorrate
attack \			
0	0.00	0.05	0.00
normal			
1	0.00	0.00	0.00
normal			
2	1.00	0.00	0.00
neptune			
3	0.01	0.00	0.01
normal			
4	0.00	0.00	0.00
normal			

	lastflag
0	20
1	15
2	19
3	21
4	21

```
[5 rows x 43 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 125973 entries, 0 to 125972
Data columns (total 43 columns):
```

#	Column	Non-Null Count	Dtype
0	duration	125973 non-null	int64

1	protocoltype	125973	non-null	object
2	service	125973	non-null	object
3	flag	125973	non-null	object
4	srcbytes	125973	non-null	int64
5	dstbytes	125973	non-null	int64
6	land	125973	non-null	int64
7	wrongfragment	125973	non-null	int64
8	urgent	125973	non-null	int64
9	hot	125973	non-null	int64
10	numfailedlogins	125973	non-null	int64
11	loggedin	125973	non-null	int64
12	numcompromised	125973	non-null	int64
13	rootshell	125973	non-null	int64
14	suattempted	125973	non-null	int64
15	numroot	125973	non-null	int64
16	numfilecreations	125973	non-null	int64
17	numshells	125973	non-null	int64
18	numaccessfiles	125973	non-null	int64
19	numoutboundcmds	125973	non-null	int64
20	ishostlogin	125973	non-null	int64
21	isguestlogin	125973	non-null	int64
22	count	125973	non-null	int64
23	srvcount	125973	non-null	int64
24	serrorrate	125973	non-null	float64
25	srvserrorrate	125973	non-null	float64
26	rerrorrate	125973	non-null	float64
27	srvrerrorrate	125973	non-null	float64
28	samesrvrate	125973	non-null	float64
29	diffsrvrate	125973	non-null	float64
30	srvdiffhostrate	125973	non-null	float64
31	dsthcount	125973	non-null	int64
32	dsthhostsrvcount	125973	non-null	int64
33	dsthhostsamesrvrate	125973	non-null	float64
34	dsthhostdiffsrvrate	125973	non-null	float64
35	dsthhostsamesrcportrate	125973	non-null	float64
36	dsthhostsrvdiffhostrate	125973	non-null	float64
37	dsthhostsererrorrate	125973	non-null	float64
38	dsthhostsrvsererrorrate	125973	non-null	float64
39	dsthhostrererrorrate	125973	non-null	float64
40	dsthhostsrvrererrorrate	125973	non-null	float64
41	attack	125973	non-null	object
42	lastflag	125973	non-null	int64

dtypes: float64(15), int64(24), object(4)
memory usage: 41.3+ MB
None

```
In [26]: # Display missing values before handling
print("Missing values before handling:")
print(df.isnull().sum())
```

Missing values before handling:

duration	0
protocoltype	0
service	0
flag	0
srcbytes	0
dstbytes	0
land	0
wrongfragment	0
urgent	0
hot	0
numfailedlogins	0
loggedin	0
numcompromised	0
rootshell	0
suattempted	0
numroot	0
numfilecreations	0
numshells	0
numaccessfiles	0
numoutboundcmds	0
ishostlogin	0
isguestlogin	0
count	0
srvcount	0
serrorrate	0
srverrorrate	0
rerrorrate	0
srvrerrorrate	0
samesrvrate	0
diffsrvrate	0
srvidfhostrate	0
dsthostcount	0
dsthostsrvcount	0
dsthostsamesrvrate	0
dsthostdiffsrvrate	0
dsthostsamesrcportrate	0
dsthostsrvidfhostrate	0
dsthosterrorrate	0
dsthostsrverrorrate	0
dsthosterrorrate	0
dsthostsrvrerrorrate	0
attack	0
lastflag	0
dtype: int64	

```
In [27]: # Separate numerical and categorical columns
numerical_columns = df.select_dtypes(include=['float64', 'int64']).
categorical_columns = df.select_dtypes(include=['object']).columns

# Replace missing values in numerical columns with the median
for col in numerical_columns:
    df[col].fillna(df[col].median(), inplace=True)

# Replace missing values in categorical columns with the most frequ
for col in categorical_columns:
    df[col].fillna(df[col].mode()[0], inplace=True)
```

```
In [28]: # Display missing values after handling
print("\nMissing values after handling:")
print(df.isnull().sum())
```

Missing values after handling:

duration	0
protocoltype	0
service	0
flag	0
srcbytes	0
dstbytes	0
land	0
wrongfragment	0
urgent	0
hot	0
numfailedlogins	0
loggedin	0
numcompromised	0
rootshell	0
suattempted	0
numroot	0
numfilecreations	0
numshells	0
numaccessfiles	0
numoutboundcmds	0
ishostlogin	0
isguestlogin	0
count	0
srvcount	0
serrorrate	0
srverrorrate	0
rerrorrate	0
srvrerrorrate	0
samesrvrate	0
diffsrvrate	0
srvdiffhostrate	0
dsthostcount	0
dsthostsrvcount	0
dsthostsamesrvrate	0
dsthostdiffsrvrate	0
dsthostsamesrcportrate	0
dsthostsrvdiffhostrate	0
dsthosterrorrate	0
dsthostsrverrorrate	0
dsthostrerrorrate	0
dsthostsrvrerrorrate	0
attack	0
lastflag	0

dtype: int64

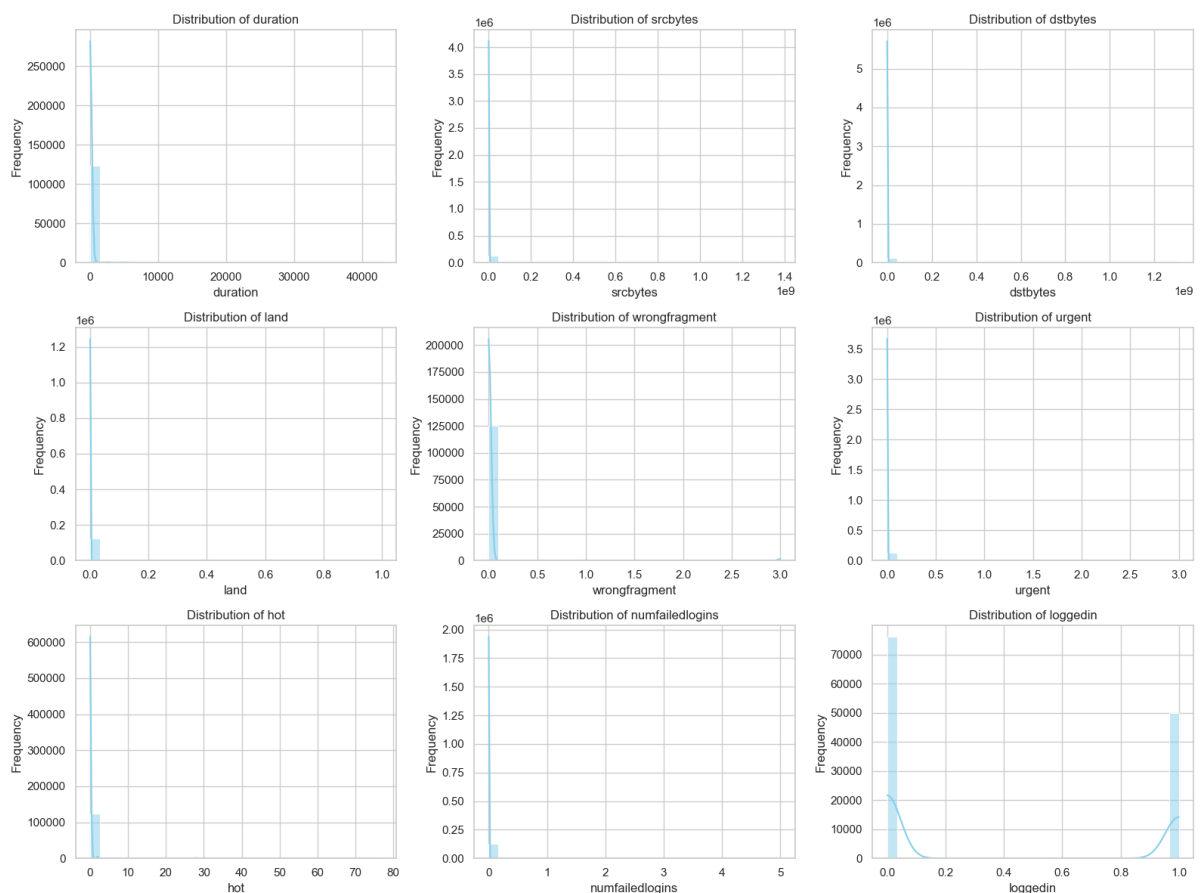
```
In [29]: # Set plot style
sns.set(style="whitegrid")

# Function to plot the distribution of numeric features
def plot_feature_distributions(data, feature_columns):
    n_cols = 3
    n_rows = (len(feature_columns) + n_cols - 1) // n_cols
    plt.figure(figsize=(16, n_rows * 4))

    for i, feature in enumerate(feature_columns, 1):
        plt.subplot(n_rows, n_cols, i)
        sns.histplot(data[feature], kde=True, bins=30, color="skyblue")
        plt.title(f"Distribution of {feature}")
        plt.xlabel(feature)
        plt.ylabel("Frequency")

    plt.tight_layout()
    plt.show()

# Select numeric columns for distribution analysis
numeric_columns = df.select_dtypes(include=['int64', 'float64']).columns
plot_feature_distributions(df, numeric_columns[:9]) # Plot for the
```

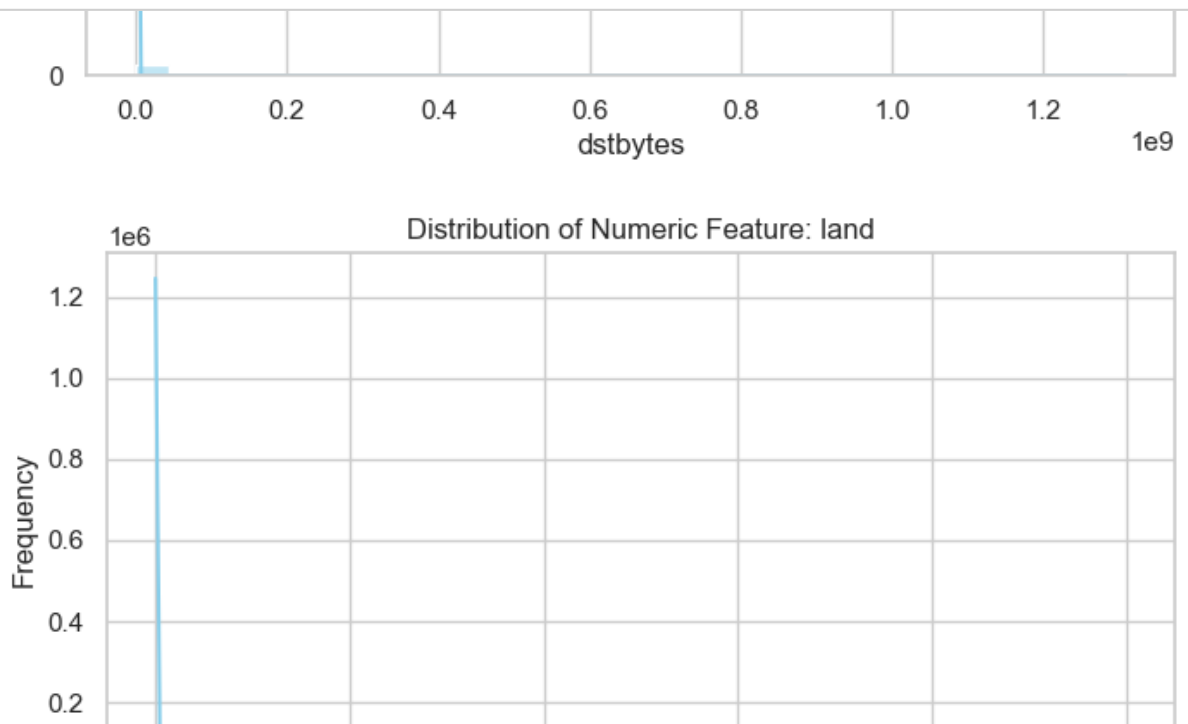


Observations:

Features like duration, srcbytes, and dstbytes have highly skewed distributions, likely influenced by extreme outliers or infrequent high values. Binary features such as land and urgent show a discrete distribution. Some features, like wrongfragment, have a significant number of zero entries, indicating sparsity.

```
In [30]: def plot_distributions(data):  
    # Separate numeric and categorical columns  
    numeric_columns = data.select_dtypes(include=['int64', 'float64'])  
    categorical_columns = data.select_dtypes(include=['object', 'category'])  
  
    # Plot distributions for numeric features  
    for column in numeric_columns:  
        plt.figure(figsize=(8, 4))  
        sns.histplot(data[column], kde=True, bins=30, color="skyblue")  
        plt.title(f"Distribution of Numeric Feature: {column}")  
        plt.xlabel(column)  
        plt.ylabel("Frequency")  
        plt.show()  
  
    # Plot distributions for categorical features  
    for column in categorical_columns:  
        plt.figure(figsize=(8, 4))  
        sns.countplot(data=data, x=column, palette="viridis")  
        plt.title(f"Distribution of Categorical Feature: {column}")  
        plt.xlabel(column)  
        plt.ylabel("Count")  
        plt.xticks(rotation=45)  
        plt.show()
```


In [31]: *# Call the function to visualize all feature distributions*
`plot_distributions(df)`



Corelation


To identify highly correlated features in your dataset and drop the ones that are redundant, we can calculate the correlation matrix and use a threshold to decide which features to drop.

In [32]: `def correlation_analysis(data):
Compute the correlation matrix
Identify numerical columns to scale/normalize
 numerical_columns = data.select_dtypes(include=['float64', 'int64'])
 corr_matrix = data[numerical_columns].corr()

Plot the heatmap
 plt.figure(figsize=(12, 10))
 sns.heatmap(corr_matrix, annot=False, cmap="coolwarm", fmt='.2f')
 plt.title("Correlation Heatmap")
 plt.show()

Return the correlation matrix for further analysis
 return corr_matrix`

In [33]: *# Call the function for correlation analysis*
 correlation_matrix = correlation_analysis(df)



In [34]: *# Check for duplicates*
 print(f"Number of duplicates before removal: {df.duplicated().sum()}")
Remove duplicates
 df_cleaned = df.drop_duplicates()
Verify if duplicates are removed
 print(f"Number of duplicates after removal: {df_cleaned.duplicated().sum()}")

Number of duplicates before removal: 0

Number of duplicates after removal: 0

```
In [35]: categorical_columns = ['protocoltype', 'service', 'flag']

# Dictionary to store mappings
label_encoders = {}
label_mappings = {}

# Apply Label Encoding and store mappings
for col in categorical_columns:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le
    label_mappings[col] = {index: label for index, label in enumerate(label_encoders[col].classes_)}

# Print the mappings for each column
for col, mapping in label_mappings.items():
    print(f"Mapping for {col}:")
    for encoded, original in mapping.items():
        print(f"    {encoded} -> {original}")
    print()

# Display the first few rows of the dataset
print("\nEncoded Dataset:")
print(df.head())

27 -> http_8001
28 -> imap4
29 -> iso_tsap
30 -> klogin
31 -> kshell
32 -> ldap
33 -> link
34 -> login
35 -> mtp
36 -> name
37 -> netbios_dgm
38 -> netbios_ns
39 -> netbios_ssn
40 -> netstat
41 -> nnsp
42 -> nntp
43 -> ntp_u
44 -> other
45 -> pm_dump
46 -> pop_2
```

```
In [36]: # Identify numerical columns to scale/normalize
numerical_columns = df.select_dtypes(include=['float64', 'int64']).

# Standardization: Mean = 0, Std Dev = 1
standard_scaler = StandardScaler()
df_standardized = df.copy()
df_standardized[numerical_columns] = standard_scaler.fit_transform(df_standardized[numerical_columns])
```

```
# Normalization: Scale to range [0, 1]
minmax_scaler = MinMaxScaler()
df_normalized = df.copy()
df_normalized[numerical_columns] = minmax_scaler.fit_transform(df[n

# Display the transformed datasets
print("Standardized Dataset (first 5 rows):")
print(df_standardized.head())

print("\nNormalized Dataset (first 5 rows):")
print(df_normalized.head())
```

```
Standardized Dataset (first 5 rows):
   duration  protocoltype  service      flag  srcbytes  dstbytes
land \
0 -0.110249    -0.124706 -0.686785  0.751111 -0.007679 -0.004919
-0.014089
1 -0.110249     2.219312  0.781428  0.751111 -0.007737 -0.004919
-0.014089
2 -0.110249    -0.124706  1.087305 -0.736235 -0.007762 -0.004919
-0.014089
3 -0.110249    -0.124706 -0.442083  0.751111 -0.007723 -0.002891
-0.014089
4 -0.110249    -0.124706 -0.442083  0.751111 -0.007728 -0.004814
-0.014089

   wrongfragment    urgent      hot  ...  dsthostsamesrvrate  \
0      -0.089486 -0.007736 -0.095076  ...              -0.782367
1      -0.089486 -0.007736 -0.095076  ...              -1.161030
2      -0.089486 -0.007736 -0.095076  ...              -0.938287
3      -0.089486 -0.007736 -0.095076  ...               1.066401
4      -0.089486 -0.007736 -0.095076  ...               1.066401

   dsthostdiffsrvrate  dsthostsamesrcportrate  dsthostsrvdiffhostr
ate \
0          -0.280282                0.069972              -0.289
103
1           2.736852                2.367737              -0.289
103
2          -0.174417               -0.480197              -0.289
103
3          -0.439078               -0.383108               0.066
252
4          -0.439078               -0.480197              -0.289
103

   dsthosterrorrate  dsthostsrverrorrate  dsthostrrerrorrate  \
0          -0.639532             -0.624871             -0.224532
1          -0.639532             -0.624871             -0.387635
2           1.608759             1.618955             -0.387635
3          -0.572083             -0.602433             -0.387635
4          -0.639532             -0.624871             -0.387635

   dsthostsrvrrerrorrate  attack  lastflag
```

```

0          -0.376387  normal  0.216426
1          -0.376387  normal -1.965556
2          -0.376387  neptune -0.219970
3          -0.345084  normal  0.652823
4          -0.376387  normal  0.652823

```

[5 rows x 43 columns]

Normalized Dataset (first 5 rows):

```

duration  protocoltype  service  flag  srcbytes  dstbytes
0          0.0          0.5  0.289855  0.9  3.558064e-07  0.000000e
+00  0.0
1          0.0          1.0  0.637681  0.9  1.057999e-07  0.000000e
+00  0.0
2          0.0          0.5  0.710145  0.5  0.000000e+00  0.000000e
+00  0.0
3          0.0          0.5  0.347826  0.9  1.681203e-07  6.223962e
-06  0.0
4          0.0          0.5  0.347826  0.9  1.442067e-07  3.206260e
-07  0.0

```

```

wrongfragment  urgent  hot  ...  dsthostsamesrvrate  dsthostdif
fsrvrate \
0          0.0          0.0  0.0  ...          0.17
0.03
1          0.0          0.0  0.0  ...          0.00
0.60
2          0.0          0.0  0.0  ...          0.10
0.05
3          0.0          0.0  0.0  ...          1.00
0.00
4          0.0          0.0  0.0  ...          1.00
0.00

```

```

dsthostsamesrcportrate  dsthostsrvdiffhostrate  dsthostserverrate
te \
0          0.17          0.00          0.
00
1          0.88          0.00          0.
00
2          0.00          0.00          1.
00
3          0.03          0.04          0.
03
4          0.00          0.00          0.
00

```

```

dsthostsrverrorrate  dsthosterrorrate  dsthostsvrerrorrate
attack \
0          0.00          0.05          0.00
normal
1          0.00          0.00          0.00

```

normal			
2	1.00	0.00	0.00
neptune			
3	0.01	0.00	0.01
normal			
4	0.00	0.00	0.00
normal			

```

    lastflag
0  0.952381
1  0.714286
2  0.904762
3  1.000000
4  1.000000

```

[5 rows x 43 columns]

Explanation of Changes:

Dropping Only One Feature from Each Pair:

For each correlated pair, only the first feature (i.e., pair[0]) is added to the correlated_features set, ensuring that only one feature from each correlated pair is dropped.

Set Data Structure for Features to Drop:

A set is used to ensure that each feature is only added once, even if it appears in multiple correlated pairs.

```

In [37]: # Select only numeric fields
numeric_df = df.select_dtypes(include=[np.number])

# Calculate the correlation matrix
correlation_matrix = numeric_df.corr()

# Set a threshold for correlation (e.g., 0.9)
threshold = 0.9

# Initialize a list to store correlated column pairs
correlated_pairs = []

# Find highly correlated features
for i in range(len(correlation_matrix.columns)):
    for j in range(i):
        if abs(correlation_matrix.iloc[i, j]) > threshold: # Check
            colname1 = correlation_matrix.columns[i]
            colname2 = correlation_matrix.columns[j]
            correlated_pairs.append((colname1, colname2))

# Print correlated column pairs

```

```

if correlated_pairs:
    print("Highly correlated column pairs (correlation > 0.9):")
    for pair in correlated_pairs:
        print(f"{pair[0]} and {pair[1]}")
else:
    print("No highly correlated column pairs found.")

# Initialize a set to keep track of features to drop
correlated_features = set()

# Keep only the first feature of each correlated pair (drop the second)
for pair in correlated_pairs:
    correlated_features.add(pair[0]) # Add only the first feature

# Drop the selected features from the original dataframe
df = df.drop(columns=correlated_features)

# Output the dropped features
print(f"\nDropped features due to high correlation: {correlated_features}")

```

Highly correlated column pairs (correlation > 0.9):

numroot and numcompromised
 srvserrorrate and serrorrate
 srvrerrorrate and rerrorrate
 dsthosterrorrate and serrorrate
 dsthosterrorrate and srvserrorrate
 dsthostsrvserrorrate and serrorrate
 dsthostsrvserrorrate and srvserrorrate
 dsthostsrvserrorrate and dsthosterrorrate
 dsthostrerrorrate and rerrorrate
 dsthostrerrorrate and srvrerrorrate
 dsthostsrvrerrorrate and rerrorrate
 dsthostsrvrerrorrate and srvrerrorrate
 dsthostsrvrerrorrate and dsthostrerrorrate

Dropped features due to high correlation: {'numroot', 'dsthosterrorrate', 'srvserrorrate', 'dsthostrerrorrate', 'dsthostsrvrerrorrate', 'srvrerrorrate', 'dsthostsrvserrorrate'}

In [38]: `df.head()`

Out[38]:

	duration	protocoltype	service	flag	srcbytes	dstbytes	land	wrongfragment	urgent	
0	0	1	20	9	491	0	0	0	0	
1	0	2	44	9	146	0	0	0	0	
2	0	1	49	5	0	0	0	0	0	
3	0	1	24	9	232	8153	0	0	0	
4	0	1	24	9	199	420	0	0	0	

5 rows × 36 columns

Feature Engineering Steps

Interaction Features: Combine numerical features to create interaction terms.

Aggregated Features: Create summary statistics like the mean, sum, or count of certain groups of features.

Polynomial Features: Introduce non-linear relationships between features by applying polynomial transformation.

Let's focus on feature engineering by combining features in a few creative ways.

```
In [39]: # Creating Interaction Features (combining numerical features)
df['src_dst_bytes_interaction'] = df['srcbytes'] * df['dstbytes']
df['num_failed_logins_hot_interaction'] = df['numfailedlogins'] * df['hot']
df['num_compromised_su_interaction'] = df['numcompromised'] * df['su']

# Aggregated Features: Summary statistics over groups of features
df['total_data_transfer'] = df['srcbytes'] + df['dstbytes'] # Total data transfer
df['total_access_operations'] = df['numfilecreations'] + df['numshells']

# Encode the 'attack' column as binary: 'normal' = 0, others = 1
df['attack_binary'] = df['attack'].apply(lambda x: 0 if x == 'normal' else 1)

# Drop any features that you may not need
df = df.drop(columns=['srcbytes', 'dstbytes', 'attack']) # Dropping original features
```



```
In [40]: df.head()
```

```
Out[40]:
```

	duration	protocoltype	service	flag	land	wrongfragment	urgent	hot	numfailedlogins
0	0	1	20	9	0	0	0	0	0
1	0	2	44	9	0	0	0	0	0
2	0	1	49	5	0	0	0	0	0
3	0	1	24	9	0	0	0	0	0
4	0	1	24	9	0	0	0	0	0

5 rows × 39 columns

Key Feature Engineering Techniques Applied:

Interaction Features:

src_dst_bytes_interaction: Multiplying source and destination bytes.

num_failed_logins_hot_interaction: Multiplying failed login attempts and the 'hot' indicator.

num_compromised_su_interaction: Multiplying the number of compromised conditions and su attempts.

Aggregated Features:

total_data_transfer: Sum of srcbytes and dstbytes.

total_access_operations: Sum of file creations, shells, and access file operations.

Polynomial Features: Polynomial transformations (degree 2) were applied to all numeric features to introduce interaction terms and squared terms, which can help capture more complex relationships between features.

Outcome: New interaction features are added, potentially revealing hidden patterns between features. Polynomial features are added, enriching the dataset with higher-order terms. The final dataset is saved as `Network_anomaly_data_feature_engineered_with_interactions.csv`.

```
In [41]: # Feature and target separation
X = df.drop(columns=['attack_binary']) # Features
y = df['attack_binary'] # Target variable

# Perform train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42, stratify=y
)

# Display the shapes of the splits
print("Training features shape:", X_train.shape)
print("Testing features shape:", X_test.shape)
print("Training target shape:", y_train.shape)
print("Testing target shape:", y_test.shape)
```

```
Training features shape: (88181, 38)
Testing features shape: (37792, 38)
Training target shape: (88181,)
Testing target shape: (37792,)
```

```
In [42]: # 3. Random Forest Classifier
rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
print(f"Random Forest Accuracy: {accuracy_score(y_test, y_pred_rf)}")
print("ROC-AUC:", roc_auc_score(y_test, rf.predict_proba(X_test)[:,
```

```
Random Forest Accuracy: 0.9994443268416596
ROC-AUC: 0.9999978162409968
```

```
In [43]: # Evaluation using classification report for better understanding o
print("\nClassification Report (Random Forest):")
print(classification_report(y_test, y_pred_rf))
```

```
Classification Report (Random Forest):
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	20203
1	1.00	1.00	1.00	17589
accuracy			1.00	37792
macro avg	1.00	1.00	1.00	37792
weighted avg	1.00	1.00	1.00	37792

```
In [44]: # Save the trained model
joblib.dump(rf, "random_forest_model.joblib")
print("Model saved successfully.")
```

```
Model saved successfully.
```

```
In [45]: # Load the model
rf_loaded = joblib.load("random_forest_model.joblib")
print("Model loaded successfully.")

# Use the loaded model for prediction
y_pred_loaded = rf_loaded.predict(X_test)
print(f"Accuracy of loaded model: {accuracy_score(y_test, y_pred_lo

Model loaded successfully.
Accuracy of loaded model: 0.9994443268416596
```

In []: